

Shift-Uniform Combes–Thomas Resolvents on Physical Gauge Cochains: A Lean-Verified Route to Exponentially Localized Covariance Roots

Lluis Eriksson*

August 1, 2026 – research preprint, version 0.2

Abstract

Let K be a coercive finite-range operator on the finite-dimensional real Hilbert space of physical lattice gauge one-cochains. A standard Combes–Thomas argument controls the inverse kernel by conjugating K with an exponential weight. Applying that argument naively to the shifted family $K + s\text{Id}$ appears to enlarge the entrywise bound by s , destroying the uniformity needed to integrate resolvents. We isolate and machine-check the missing algebraic fact: scalar identity shifts cancel exactly from the tilt defect,

$$(K + s\text{Id})_\theta - (K + s\text{Id}) = K_\theta - K.$$

Consequently, one fixed Combes–Thomas budget for K controls every $K + s\text{Id}$, $s \geq 0$. If K is coercive with constant $c > 0$, has range R , single-bond kernel bound M , and ball-cardinality bound N_R , then under $M(e^{\theta R} - 1)N_R \leq c/2$, every right inverse C_s of $K + s\text{Id}$ satisfies

$$\|(C_s)_{q,p}\|_{\text{op}} \leq \frac{2}{c+s} e^{-\theta \text{dist}(q,p)}.$$

The cancellation identity, shifted coercivity, uniform defect estimate, and resolvent kernel theorem are formalized in Lean 4 on the repository’s actual physical gauge-cochain types. At paper level, self-adjoint positivity and the Stieltjes formula then give

$$\|(K^{-1/2})_{q,p}\|_{\text{op}} \leq \frac{2}{\sqrt{c}} e^{-\theta \text{dist}(q,p)}.$$

The fractional-power decay itself is classical in matrix analysis; the contribution claimed here is the shift-uniform, source-facing, machine-checked route with explicit constants and its precise connection to an open covariance-root interface in a formalized lattice Yang–Mills programme. We state the remaining gap without compression: the Stieltjes functional calculus, the positive root, and the final root certificate have not yet been formalized, and the fixed-volume coercivity used here is not volume-uniform.

1 Claims and status discipline

This paper separates three levels that are easy to conflate.

*This manuscript and part of the formal development were produced with AI assistance under the repository’s audit protocol. Claims labelled FORMALIZED were checked by Lean 4 against a pinned Mathlib and a targeted axiom oracle. Paper-level functional-calculus arguments are labelled PROVED and are not represented as Lean theorems.

Status	Meaning in this manuscript
FORMALIZED	A named Lean theorem compiles on the pinned tree and the targeted axiom oracle reports only <code>propext</code> , <code>Classical.choice</code> , and <code>Quot.sound</code> .
PROVED	A complete finite-dimensional mathematical proof is given here, but the proof is not yet a Lean declaration.
OPEN	The statement is neither assumed silently nor claimed. It is a named frontier for subsequent formalization or analysis.

The central resolvent result, Theorem 5.2, is FORMALIZED. The Stieltjes passage to $K^{-1/2}$, Theorem 6.1, is PROVED at paper level. Construction of the positive root as a continuous linear map in the repository and discharge of the existing `PhysicalLocalizedCovarianceRootCertificate` are OPEN. No continuum limit, mass gap, or volume-uniform positive coercivity constant is claimed.

The mathematical phenomenon is not presented as new. Exponential decay of inverses of banded matrices goes back at least to Demko–Moss–Smith [2]; decay bounds for completely monotone functions, including negative fractional powers, were developed directly by Benzi–Simoncini [3]. Our novelty claim is deliberately narrower: a machine-checked shift-cancellation mechanism and a reusable constant-explicit resolvent theorem on the physical gauge-cochain interface used by the surrounding formal development.

2 Physical cochains and localization predicates

Fix $d, N, N_c \in \mathbb{N}$ with $N > 0$. Let $\Lambda_N = (\mathbb{Z}/N\mathbb{Z})^d$ and let $\mathcal{B}_{d,N}$ be the finite set of positively oriented lattice bonds. The internal coordinate space is the real coordinate model of $\mathfrak{su}(N_c)$,

$$V_{N_c} \cong \mathbb{R}^{N_c^2-1}.$$

The finite-dimensional real Hilbert space of physical one-cochains is

$$\mathcal{H}_{d,N,N_c} := \ell^2(\mathcal{B}_{d,N}; V_{N_c}).$$

In Lean these are, respectively, `PhysicalBond d N`, `SUNLieCoord Nc`, and `PhysicalGaugeOneCochain d N Nc`.

For $p \in \mathcal{B}_{d,N}$ and $v \in V_{N_c}$, write $\delta_p v$ for the cochain supported at p with value v . This is `singlePhysicalBondCochain`. Let

$$\text{dist} : \mathcal{B}_{d,N} \times \mathcal{B}_{d,N} \longrightarrow \mathbb{N}$$

be symmetric, satisfy the triangle inequality, and obey $\text{dist}(p, p) = 0$.

Let $K : \mathcal{H} \rightarrow \mathcal{H}$ be a continuous linear map. The formal finite-range and kernel hypotheses have the following mathematical content:

$$\text{dist}(q, p) > R \implies (K\delta_p v)(q) = 0, \tag{1}$$

$$\|(K\delta_p v)(q)\| \leq M\|v\|. \tag{2}$$

The cardinality parameter N_R controls every metric ball,

$$\#\{y : \text{dist}(x, y) \leq R\} \leq N_R. \tag{3}$$

Finally, K is coercive with constant c if

$$c\|f\|^2 \leq \langle f, Kf \rangle \quad (f \in \mathcal{H}). \quad (4)$$

The corresponding Lean predicate is `IsCoerciveCLM K c`.

3 Exponential conjugation

Fix a root bond r and $\theta \in \mathbb{R}$. Define the diagonal tilt

$$(T_{\theta,r}f)(q) = e^{\theta \text{dist}(r,q)} f(q) \quad (5)$$

and the conjugated operator

$$K_{\theta,r} = T_{\theta,r} K T_{-\theta,r}. \quad (6)$$

Because the bond set is finite, these are bounded maps without a domain issue. The formal development proves $T_{\theta,r} T_{-\theta,r} = \text{Id}$ and, on a single-bond probe,

$$(K_{\theta,r} \delta_p v)(q) = e^{\theta(\text{dist}(r,q) - \text{dist}(r,p))} (K \delta_p v)(q). \quad (7)$$

On the range of K , the triangle inequality gives $|\text{dist}(r,q) - \text{dist}(r,p)| \leq R$. The block Schur estimate in the formal library then yields

$$\|K_{\theta,r} - K\| \leq \Delta_\theta, \quad \Delta_\theta := M(e^{\theta R} - 1)N_R, \quad \theta \geq 0. \quad (8)$$

This is the existing Combes–Thomas engine. The issue for a Stieltjes construction is whether the same Δ_θ controls all positive scalar shifts.

4 The scalar-shift cancellation

Lemma 4.1 (identity is fixed; FORMALIZED). *For every r , θ , and physical gauge-cochain space,*

$$T_{\theta,r} \text{Id} T_{-\theta,r} = \text{Id}.$$

In Lean this is `physicalTiltConjCLM_id`.

Proof. It is the inverse relation $T_{\theta,r} T_{-\theta,r} = \text{Id}$. The Lean proof unfolds the conjugation and rewrites with the already formalized composition identity. $\square$

Lemma 4.2 (shift cancellation; FORMALIZED). *For $s \in \mathbb{R}$,*

$$(K + s\text{Id})_{\theta,r} - (K + s\text{Id}) = K_{\theta,r} - K. \quad (9)$$

In Lean this is `physicalTiltConjCLM_add_smul_id_sub`.

Proof. Conjugation is linear in the operator. Lemma 4.1 gives $(K + s\text{Id})_{\theta,r} = K_{\theta,r} + s\text{Id}$; subtracting $K + s\text{Id}$ proves (9). The Lean proof separately establishes additivity and scalar compatibility of conjugation and closes the operator identity by abelian-group normalization. $\square$

The importance of (9) is not its algebraic difficulty but its location in the estimate. A naive application of the unshifted theorem to $K + s\text{Id}$ asks for a kernel bound on the whole shifted operator and therefore introduces a spurious term proportional to s . Equation (9) shows that the tilt only sees the off-diagonal commutator; the scalar diagonal shift is invisible.

Corollary 4.3 (uniform defect; FORMALIZED). *Under (1)–(3), for every $s \in \mathbb{R}$,*

$$\|(K + s\text{Id})_{\theta,r} - (K + s\text{Id})\| \leq \Delta_\theta.$$

This is `norm_physicalTiltConj_add_smul_id_sub_le`.

5 Shift-uniform coercivity and resolvents

The scalar shift also improves coercivity exactly.

Lemma 5.1 (shifted coercivity; FORMALIZED). *If (4) holds, then*

$$(c + s)\|f\|^2 \leq \langle f, (K + s\text{Id})f \rangle$$

for every real s . In Lean this is `isCoerciveCLM_add_smul_id`.

Combining Lemma 5.1 with Corollary 4.3 gives the shift-uniform CT2 endpoint

$$\langle f, (K + s\text{Id})_{\theta,r} f \rangle \geq (c + s - \Delta_\theta)\|f\|^2. \quad (10)$$

This is `isCoerciveCLM_physicalTiltConj_add_smul_id`. If $s \geq 0$ and $\Delta_\theta \leq c/2$, then

$$c + s - \Delta_\theta \geq \frac{c + s}{2}. \quad (11)$$

The corresponding formal theorem is `isCoerciveCLM_physicalTiltConj_add_smul_id_half`.

Theorem 5.2 (shift-uniform Combes–Thomas resolvent; FORMALIZED). *Assume $c > 0$, $M \geq 0$, $\theta > 0$, equations (1)–(4), and*

$$M(e^{\theta R} - 1)N_R \leq \frac{c}{2}. \quad (12)$$

For $s \geq 0$, let C_s be any right inverse of $K + s\text{Id}$:

$$(K + s\text{Id})C_s = \text{Id}.$$

Then, for all bonds p, q and $v \in V_{N_c}$,

$$\|(C_s \delta_p v)(q)\| \leq \frac{2}{c + s} e^{-\theta \text{dist}(q,p)} \|v\|. \quad (13)$$

This is `physicalCovariance_exponentialKernelBound_of_coercive_add_smul_id`.

Proof architecture. Root the tilt at the source bond p and set $y = T_{\theta,p} C_s \delta_p v$. Since $\text{dist}(p, p) = 0$, the source probe is fixed by its own tilt. The right-inverse relation and exact untilting give

$$(K + s\text{Id})_{\theta,p} y = \delta_p v.$$

Coercivity at $(c + s)/2$ and Cauchy–Schwarz imply

$$\|y\| \leq \frac{2}{c + s} \|v\|.$$

At the target bond,

$$(C_s \delta_p v)(q) = e^{-\theta \text{dist}(p,q)} y(q).$$

The coordinate projection has norm at most one, and symmetry of the distance gives (13). Every equality and inequality in this proof architecture is represented in the cited Lean theorem. $\square$

Remark 5.3 (why the denominator is integrable). *The key improvement over a shift-blind estimate is not merely that the decay rate remains positive. The amplitude retains the spectral shift as $(c + s)^{-1}$. Substituting $s = t^2$ produces an integrable majorant on $[0, \infty)$, which is exactly what the inverse-square-root Stieltjes formula requires.*

6 The covariance-root corollary

The next theorem is ordinary finite-dimensional spectral calculus. It is included with a complete proof to identify the exact remaining Lean task.

Theorem 6.1 (localized inverse square root; PROVED). *In addition to the hypotheses of Theorem 5.2, assume that K is self-adjoint and positive with $\sigma(K) \subset [c, \infty)$, $c > 0$. Then the unique positive inverse square root $K^{-1/2}$ satisfies*

$$\|(K^{-1/2} \delta_p v)(q)\| \leq \frac{2}{\sqrt{c}} e^{-\theta \text{dist}(q,p)} \|v\|. \quad (14)$$

Moreover,

$$\|K^{-1/2}\| \leq c^{-1/2}, \quad K^{-1/2} K^{-1/2} = K^{-1}.$$

Proof. The spectral theorem gives the norm-convergent Stieltjes identity

$$K^{-1/2} = \frac{2}{\pi} \int_0^\infty (K + t^2 \text{Id})^{-1} dt. \quad (15)$$

Indeed, after diagonalizing K , the operator identity reduces to the scalar identity, valid for $\lambda > 0$,

$$\frac{2}{\pi} \int_0^\infty \frac{dt}{\lambda + t^2} = \frac{2}{\pi} \frac{\pi}{2\sqrt{\lambda}} = \lambda^{-1/2}. \quad (16)$$

There is no interchange difficulty in finite dimension. Apply Theorem 5.2 with $s = t^2$ and integrate the pointwise norm majorant:

$$\begin{aligned} \|(K^{-1/2} \delta_p v)(q)\| &\leq \frac{2}{\pi} \int_0^\infty \frac{2}{c + t^2} e^{-\theta \text{dist}(q,p)} \|v\| dt \\ &= \frac{2}{\sqrt{c}} e^{-\theta \text{dist}(q,p)} \|v\|. \end{aligned}$$

The spectral bounds and the square identity follow by applying $\lambda \mapsto \lambda^{-1/2}$ on $\sigma(K)$. $\square$

Remark 6.2 (constant accounting). *The spectral norm of $K^{-1/2}$ is at most $c^{-1/2}$, whereas the kernel amplitude obtained from the half-budget CT argument is $2c^{-1/2}$. The factor two is the price of reserving half the coercivity uniformly for every tilt. We make no optimality claim for this constant or for the admissible rate θ .*

7 Connection to the formal covariance-root frontier

The repository already contains the proposition-valued structure `PhysicalLocalizedCovarianceRootCertificate`. Given a precision K , a covariance C , and a proposed root B , it requires:

1. a localized covariance certificate for (K, C) ;
2. $B \circ B = C$;
3. a norm bound on B ;
4. self-adjointness of B in inner-product form;
5. positivity of B ;
6. a single-bond kernel bound on B .

Theorem 6.1 supplies the mathematical content of items 2–6 when $C = K^{-1}$ and $B = K^{-1/2}$, while Theorem 5.2 supplies the uniform resolvent estimate needed for item 6. It does not yet supply a Lean term B .

The local-SPD route in `PhysicalGaugeLocalSPDPrecisionRoot.lean` is complementary. It develops Neumann and inverse-square-root coefficient majorants under a normalized small-perturbation hypothesis $K = a(\text{Id} - L)$ with $\|L\| < 1$, but deliberately leaves the continuous-linear root and its kernel bound as source fields. The present route replaces that smallness hypothesis by coercivity plus finite range. Neither route currently closes the positive-root object in Lean.

The exact next formal chain is therefore

$$\begin{array}{l}
 \text{self-adjoint coercive } K \longrightarrow \text{positive functional-calculus root } B = K^{-1/2} \\
 \longrightarrow \text{Stieltjes identity for } B \longrightarrow \text{integrated single-bond kernel bound} \\
 \longrightarrow \text{PhysicalLocalizedCovarianceRootCertificate.}
 \end{array} \tag{17}$$

Mathlib already provides a continuous functional calculus and positive square roots in suitable star-algebra settings. The nontrivial engineering question is to bridge the repository’s real continuous-linear-map representation to a functional-calculus carrier while preserving the single-bond kernel API.

Post-review implementation audit. The available library route is not a one-line application of an existing operator square-root theorem. Mathlib exposes positive square roots for positive semidefinite real matrices through continuous functional calculus, whereas the direct continuous-functional-calculus infrastructure for continuous linear endomorphisms is presently available in the complex Hilbert case, not as the required real-operator instance. A non-circular Lean closure therefore has four explicit obligations: choose an orthonormal coordinate model for the physical cochain space; transport the positive matrix root back to a real continuous linear map; identify that root with the norm-convergent Stieltjes integral; and pass the integrated coefficient estimate through the single-bond kernel predicate. Supplying the last predicate as an assumption would merely restate the open certificate and is deliberately rejected here.

8 Prior art and exact delta

Combes and Thomas introduced exponential conjugation in spectral analysis [1]. Demko, Moss, and Smith proved exponential decay of inverse band matrices with spectral-rate information [2]. Benzi and Simoncini treated completely monotone functions of Hermitian banded and sparse matrices, explicitly including negative fractional powers [3]. Thus neither Theorem 6.1 nor the use of a resolvent integral is claimed as a new analytic phenomenon.

Antecedent	Overlap	Exact difference claimed here
Combes–Thomas [1]	Exponential weighting and resolvent localization	Finite physical gauge-cochain API, explicit Schur budget, exact scalar-shift cancellation, and machine-checked source-rooted extraction.
Demko–Moss–Smith [2]	Off-diagonal decay of inverses	Metric finite-range block operators rather than scalar band matrices; the main delta is formal verification, not a sharper rate.
Benzi–Simoncini [3]	Decay of negative fractional powers via matrix-function representations	No claim of analytic priority. We provide a Lean-verified resolvent family with the precise $(c+s)^{-1}$ amplitude consumed by the physical covariance-root interface.
Repository fixed-volume CT chain [8]	Formal inverse decay with amplitude $2/c$	The old theorem did not preserve a scalar shift; the new theorem is uniform for all $s \geq 0$ and enables the Stieltjes integral.
Repository local-SPD root frontier	Inverse-square-root coefficient and kernel majorants	Present hypotheses are coercivity plus finite range, not $\ L\ < 1$; the root object remains open in both lanes.

The literature search for this manuscript was performed on August 1, 2026, using combinations of “Combes–Thomas”, “inverse square root”, “fractional powers”, “sparse matrix functions”, “Stieltjes”, and “Lean”. We found no machine-checked Combes–Thomas inverse-square-root theorem, but this negative search result is not an exhaustiveness claim. The paper’s defensible dominance axis is reproducibility and integration into a concrete formal physics interface, not classical analytic novelty.

9 Adversarial audit and limitations

Attack 1: the theorem is already known under matrix-function language. Correct at the analytic level. This is why the title says “Lean-verified route” and not “new localization theorem”, and why Section 8 treats negative fractional powers as direct prior art.

Attack 2: the result is only finite volume. Correct. The formal physical cochain space is finite-dimensional. A separate machine-checked Poincaré-wall result shows that the current unscaled block-Poincaré route does not provide a positive coercivity constant uniform in volume for the relevant dimensions [8]. The present theorem preserves whatever c is supplied; it does not manufacture a uniform c .

Attack 3: coercivity does not imply a positive square root for a non-self-adjoint operator. Correct. The formal resolvent theorem does not need self-adjointness, but Theorem 6.1 adds self-adjoint positivity explicitly. Omitting it would be a false bridge.

Attack 4: a right inverse family is not a measurable or continuous resolvent map. In finite dimension the canonical inverse $(K + t^2 \text{Id})^{-1}$ is continuous in t for a positive self-adjoint K . The Lean theorem is intentionally pointwise in an arbitrary right inverse; the continuity and canonical identification belong to the open functional-calculus layer.

Attack 5: the diagonal is not genuinely off-diagonal data. The formal global kernel bound on the base operator still includes diagonal entries. The improvement is that the *shift* is eliminated exactly from the tilt defect. We do not claim a new off-diagonal norm predicate in this version.

Attack 6: the constants may be poor. Correct. The factor two and the Schur cardinality N_R are explicit but not optimized. Polynomial or best-approximation methods may give sharper rates for banded Hermitian matrices. The purpose here is a robust formal interface.

Kill criteria. The programme should be reformulated if any of the following occurs: (i) the exact shift identity fails under the actual root representation; (ii) the functional-calculus bridge cannot preserve the single-bond kernel API without assuming the desired kernel bound; (iii) the root certificate requires a volume-uniform c from the already refuted block normalization; or (iv) a prior formalization is found with the same physical types, hypotheses, constants, and endpoint. None is currently observed, but (ii) is the real technical bottleneck.

10 Lean theorem map and reproducibility

10.1 Theorem-to-artifact map

The two source modules are `YangMills/RG/PhysicalCoerciveCombesThomas.lean` (CT1/CT2) and `YangMills/RG/PhysicalCoerciveCombesThomasInverse.lean` (CT3/CT4).

Lean declaration	File	Status
<code>physicalTiltConjCLM_id</code>	CT1/CT2 module	EXACT
<code>physicalTiltConjCLM_add_smul_id_sub</code>	CT1/CT2 module	EXACT
<code>isCoerciveCLM_add_smul_id</code>	CT1/CT2 module	EXACT
<code>norm_physicalTiltConj_add_smul_id_sub_le</code>	CT1/CT2 module	EXACT
<code>isCoerciveCLM_physicalTiltConj_add_smul_id</code>	CT1/CT2 module	EXACT
<code>isCoerciveCLM_physicalTiltConj_add_smul_id_half</code>	CT3/CT4 module	EXACT
<code>physicalCovariance_exponentialKernelBound_of_coercive_add_smul_id</code>	same	EXACT
Stieltjes identity and Theorem 6.1	this manuscript	paper
Positive CLM root and root certificate	not yet implemented	open

10.2 Snapshot and commands

Repository: <https://github.com/lluiseriksson/THE-ERIKSSON-PROGRAMME>. The formal artifact is the dedicated Git commit `b542b3f2f066b60a914977f2a50160e403a53367` on branch `codex/stieltjes-root`; its parent is `81e9cf064d6771f986978ccffb96018c78b5e459`. The commit contains exactly the two theorem modules, the oracle registrations, and the verification-ledger entry. Thus the formal source cited here is an immutable Git object rather than an uncommitted working-tree patch. Remote publication and a hosted CI run of that commit remain release actions and are not claimed by this manuscript.

Toolchain: Lean `leanprover/lean4:v4.29.0-rc6`; Mathlib revision `0764272048...` (the full revision is recorded in `lake-manifest.json`). The focused checks are:

```
lake env lean YangMills/RG/PhysicalCoerciveCombesThomas.lean
lake env lean YangMills/RG/PhysicalCoerciveCombesThomasInverse.lean
lake build +YangMills.RG.PhysicalCoerciveCombesThomas:olean
lake build +YangMills.RG.PhysicalCoerciveCombesThomasInverse:olean
```

All four completed successfully on August 1, 2026. A targeted oracle checked the seven declarations in the table; each printed exactly

`[proptext, Classical.choice, Quot.sound]`.

The declarations are also registered in `oracle_check.lean`. The development contains no `sorry` in the two modified modules.

11 Author-corpus publication provenance

The internal antecedent corpus contains public records whose corrective replacement PDFs had been submitted but not yet processed by the ai.viXra administrator as of August 1, 2026. They must not be cited as though the replacement versions were already public, and the superseded claims in the currently visible PDFs must not be used as support. All thirteen forms were reported sent; all thirteen remain **ENVIADO/PENDIENTE**; none should be resent merely because the public pages lag.

Record	Submitted target	Claim-level instruction while pending
2602.0052	v3	Restores the <i>Interface Lemmas</i> manuscript; visible v2 contains misfiled Morse–Bott material.
2602.0036	v3	Retracts the unproved Ricci/O’Neill chain and dependent corollaries.
2602.0033	v3/R30	Replaces false exact gap doubling by a correct sufficient condition; R29 was never submitted.
2602.0085	v2	Corrects the linearized kernel and restores the omitted pure-gauge null space.
2602.0084	v2	Corrects imported 0085 claims and independent invalid steps.
2602.0035	v2	Installs the corrected Morse–Bott spectral reduction with the required reservations.
2602.0038	v3	Narrows the claim to a zero-momentum propagator bound in a simplified Gribov–Zwanziger lattice measure.

Record	Submitted target	Claim-level instruction while pending
2602.0041	v4	Makes the H-XSD/H-DOB conditionality explicit in title and abstract.
2601.0115	v3	Preserves the science and repairs the clipped Table 1.
2512.0073	v2 notice	Marked superseded by 2601.0023v2, which corrects the Davies/Bohr decomposition and TFIM implementation.
2601.0047	v3 notice	Marked superseded by 2601.0051v2, which retains the ED benchmark and adds tensor-network results.
2607.0035	v2 notice	Retains the reduced-model analytic chain but points to 2607.0039v1 for the complete theorem on original edges and tree-cotree closure. It is not withdrawn, but it is not an autonomous proof of the full theorem.
2607.0023	v2 notice	Points the terminal-sign claim to 2607.0089v1; 2607.0089v1 remains REVIEW-PENDING until the exact-to-Arb/interval certification is repeated.

The controlling audit records are `docs/publication-audit/SUPERSESSSION-MATRIX.md` and `docs/publication-audit/SUBMISSION-STATUS-20260801.md` in the publication audit worktree. Public promotion requires a fresh download and comparison of version, page count, SHA-256, title, abstract, comments, and rendered pages. This manuscript uses none of the corrected-but-not-yet-public claims as a premise.

12 Conclusion

The useful observation is exact and small: scalar identity shifts disappear from exponential-conjugation defects. Its consequence is structurally larger: a single fixed Combes–Thomas budget controls an entire positive resolvent family with the integrable amplitude $(c + s)^{-1}$. That family is now a machine-checked theorem on the physical gauge-cochain types that the covariance-root interface actually consumes.

The paper-level Stieltjes integral then produces exponential localization of $K^{-1/2}$ with amplitude $2/\sqrt{c}$. This is not advertised as a new result in matrix analysis. What remains potentially valuable is the formal bridge: turning classical functional calculus into an auditable producer of a localized Gaussian covariance-root certificate for lattice gauge renormalization. Closing that bridge, while respecting the fixed-volume Poincaré wall, is the next theorem rather than a conclusion of this paper.

References

- [1] J.-M. Combes and L. Thomas, *Asymptotic behaviour of eigenfunctions for multiparticle Schrödinger operators*, Communications in Mathematical Physics **34** (1973), 251–270. [doi:10.1007/BF01646473](https://doi.org/10.1007/BF01646473).
- [2] S. Demko, W. F. Moss, and P. W. Smith, *Decay rates for inverses of band matrices*, Mathematics of Computation **43** (1984), 491–499. [doi:10.1090/S0025-5718-1984-0758197-9](https://doi.org/10.1090/S0025-5718-1984-0758197-9).

- [3] M. Benzi and V. Simoncini, *Decay bounds for functions of Hermitian matrices with banded or Kronecker structure*, SIAM Journal on Matrix Analysis and Applications **36** (2015), 1263–1282. doi:10.1137/151006159.
- [4] T. Balaban, *Propagators and renormalization transformations for lattice gauge theories. I*, Communications in Mathematical Physics **95** (1984), 17–40.
- [5] T. Balaban, *Propagators and renormalization transformations for lattice gauge theories. II*, Communications in Mathematical Physics **96** (1984), 223–250.
- [6] L. de Moura and S. Ullrich, *The Lean 4 theorem prover and programming language*, CADE-28, Lecture Notes in Computer Science 12699 (2021), 625–635.
- [7] The mathlib community, *The Lean mathematical library*, CPP 2020, 367–381.
- [8] L. Eriksson, *The volume-uniform Poincaré walls: machine-checked obstructions for flat and fluctuation-sector block-Poincaré routes to Combes–Thomas coercivity in lattice Yang–Mills*, repository preprint (2026), papers/poincare-wall/poincare_wall.tex.