

The Reconstructed Theory Has One Mass: a Machine-Checked Volume-Uniform Spectral Gap with Exact Identification Against the Gibbs Sums

Lluís Eriksson*

2026-08-05 — complete manuscript, zero live slots across all three fabrication stages;
submission is the owner's decision

Abstract

For the spatial $\mathbb{Z}_2$ (Ising-slice) system in the Dobrushin window $2 \tanh |\beta| + 2 \tanh |\gamma| \leq \alpha < 1$, we machine-check in Lean 4 an end-to-end chain: (i) the Osterwalder–Schrader (site-form) reconstruction of the transfer operator is unitarily conjugate — by the explicit $\sqrt{w}$ boundary dressing — to the symmetrised Dobrushin kernel; (ii) the *unnormalised Gibbs sums themselves* are exact matrix elements of the reconstructed operator's powers, $\text{gibbsPathSum } w \beta N A B = \lambda^N \langle \hat{T}^N Q A, Q B \rangle$ and the partition function is the same shape at the dressed constant — identities, not bounds; and (iii) there is *one* mass $m > 0$ such that for *every* spatial extent L the projected operator norm is at most e^{-m} and every *mixed* connected correlator of the reconstructed operator obeys $|\langle u, \hat{T}^n v \rangle - \langle \Omega, u \rangle \langle \Omega, v \rangle| \leq \|u\| \|v\| (e^{-m})^n$, the zero-time case included — constants per-observable, rate alone uniform in the volume. All quantifiers live in the Lean statements; the prose here copies the kernel's cut. To situate the claim in the 2026 landscape of machine-checked constructive field theory: the abstract construction (reflection-positive data $\rightarrow$ physical Hilbert space $\rightarrow$ self-adjoint contractive transfer operator) has been machine-checked, sorry-free, in the Harvard formalization programme, *with the spectral gap carried as an undischarged hypothesis* (**GappedTransfer**); the Osterwalder–Schrader axioms have been machine-checked for the free field; and measure-side Dobrushin clustering has been machine-checked for lattice Yang–Mills. What we add, and what to our knowledge no machine-checked artefact yet contains, is the combination {concrete interacting model, gap *proved* rather than hypothesised, uniform in the spatial volume} on the reconstructed operator, *identified exactly against the raw Gibbs sums* — i.e. a concrete discharge of exactly the hypothesis shape the abstract frameworks carry. We do not claim Wightman reconstruction, a continuum limit, $SU(N)$, or any progress on the Clay problem.

1 Introduction and landscape

The Osterwalder–Schrader route builds a physical theory out of Euclidean data: a reflection-positive form gives a Hilbert space, and composition with the time shift gives a self-adjoint transfer operator. For a lattice slice system the construction is elementary — the content is

*With machine-checked artefacts produced on the programme's Colab plane; toolchain `leanprover/lean4:v4.29.0-rc6`, Mathlib pinned to `0764272048...`. Repository: THE-ERIKSSON-PROGRAMME, branch `d3-closure`, module anchor commit `6ef9aee03ffaec062358b81afa8b3f89610d8ea`; verification anchor `639296c4d43219b0966e4d273eebd4222ef9396c` (fresh clone, Addendum 628).

never the construction but the *spectrum*: a spectral gap at every fixed spatial extent L is classical (Perron–Frobenius; in this programme it is the earlier `coupled_gap_all_eigenvalues`), and it degenerates as $L \rightarrow \infty$ unless something uniform survives. What a thermodynamic limit can actually consume is *one* $m > 0$ valid for every L at once.

That uniform input exists in this repository: the Dobrushin corollary `dobrushin_ising_uniform_gap` produces, inside the window $2 \tanh |\beta| + 2 \tanh |\gamma| \leq \alpha < 1$, a single $m > 0$ bounding the projected transfer operator of the tilted kernel at every extent. This paper’s contribution is the *identification*, in two stages. Stage one (Section 3): the operator that bound governs *is* — unitarily, by the explicit $\sqrt{w}$ boundary dressing, with no spectral-transport theory — the transfer operator forced on the reconstructed (site-form) theory. Stage two (Section 5): the *unnormalised Gibbs sums themselves* are exact matrix elements of that operator’s powers — λ^N times an inner product, an identity valid at every real β and every positive weight, with the window entering only when the gap is invoked. The theorems are deliberately cheap: the same witnesses as D-6, no new estimate. Cheapness is the honesty of the claim — the uniformity is real because it is D-6’s, and the reconstruction adds commuting squares, machine-checked.

1.1 Machine-checked prior work, cited before a referee does

The following table is load-bearing and its rows were collected by a four-modality search on 2026-08-04 (evidence: `docs/OS-R-PRIORITY-COLLATION.md` in the repository; the four raw agent reports are appended there verbatim).

Artefact	What is machine-checked	What is not
<code>mrdouglasny/reflection-positivity</code> (sorry-free ~2026-06)	RP data $\rightarrow$ Hilbert space (completion/quotient) $\rightarrow$ self-adjoint contraction T ; the dictionary $\langle [f], T^n[g] \rangle = \int f(g \circ \tau^n \circ \theta) d\mu$; trace formulas; susceptibility bounds <i>from</i> a gap	the gap itself: carried as the GappedTransfer <i>hypothesis</i> (GroundGap.lean); no spatial-volume-uniform statement; concrete instance is finite-lattice RP only
<code>OSforGFF</code> (arXiv:2603.15770)	all five GJ/OS axioms for the $d=4$ Gaussian free field, 0 sorries	reconstruction (listed as next milestone); any interacting model
<code>xiyin137/OSreconstruction</code>	Wightman-side GNS lane of the full $OS \leftrightarrow$ Wightman theorem	the theorem: 53 sorries + 12 axioms at last README snapshot
<code>mrdouglasny/lgt</code> (complete 2026-05-04)	Chatterjee-style Dobrushin uniqueness for lattice YM, $U(n)$, with exponential clustering whose constants depend only on (n, d, β)	an operator gap; any reconstruction; the statement is per-volume on the measure side
<code>pphi2</code>	transfer-matrix spectral machinery for φ_2^4	rests on 27–29 project axioms
This paper	the $\sqrt{w}$ -dressing unitary; the transported volume-uniform gap; the <i>exact</i> identities $\text{gibbsPathSum} = \lambda^N \langle \hat{T}^N QA, QB \rangle$ and $Z_N = \lambda^N \langle \hat{T}^N Q\mathbf{1}, Q\mathbf{1} \rangle$; volume-uniform <i>mixed</i> clustering with the zero-time case (<code>os_reconstruction_measure_uniform</code> , core 8481, oracle 3026, 0 sorry — measured, Addendum 622)	Wightman axioms; continuum limit; $SU(N)$; Clay; a volume-uniform CONSTANT for the normalised bound (the rate is uniform in L ; $C(A, B)$ is not); the N -uniform denominator floor IS a theorem (Section 8, 9)

There is a coordinated, active programme (Douglas et al., Harvard CMSA) whose declared live fronts are the OS reconstruction theorem, $P(\varphi)_2$, and YM mass gap at strong coupling, and whose framework carries the spectral gap as an undischarged hypothesis. The endpoint below is a concrete discharge of that hypothesis *shape* (different framework, same mathematical hole), which is why the comparison table above is printed in full before any claim wording.

2 The two geometries and the forced operator

Everything in this section is prior work of the repository, cited by module (the public trace of this lane is the ai.viXra 2607.009x cluster); no new mathematics appears here. The slice system carries two pairings on complex boundary vectors: the *site* form, which weighs by $1/w$ (`siteForm`, module `SpatialReconstruction`), and the *bond* form carrying the decoupled kernel $\text{spatialKernel } \beta$ (`bondForm`). The transfer operator is *forced*: $T = D_w K$ is the unique operator with $\text{site}(u, Tv) = \text{bond}(u, v)$ (`transferOp_unique`), self-adjoint for the site form at every β

(`transferOp_selfAdjoint`) and positive for $\beta \geq 0$; reflection positivity of the underlying measure is the lane’s earlier bond/site pair (modules `SpatialReflection`, `SpatialOS`). On the measure side, the same repository carries the Gibbs bridge (module `SpatialGibbs`): the unnormalised two-point path sum `gibbsPathSum` and partition function `gibbsPartition` of the coupled slice system, with the in-tree operator bridge `gibbsPathSum_eq_iterate` (the path sum as an N -fold kernel iterate against dressed observables, $\text{dress } w A = \sqrt{w} A$). On the Dobrushin side it carries `symWeighted` $= \sqrt{w} K \sqrt{w}$ (module `PerronGap`), its normalisation `tiltKernel` $= \text{symWeighted} / \lambda$ (module `DobrushinTilt`), the Euclidean packaging `opOf`, `vacOf`, `vacuumTransfer_opOf` (module `DobrushinTransport`), and the corollary `dobrushin_ising_uniform_gap` (module `DobrushinCorollary`). The seam fact this paper depends on — one `spatialKernel` declaration serving both geometries — is enforced by the elaborator in the module of Section 3.

2.1 The OS provenance: reflection positivity, collapse, null space, physical quotient

The operator of Section 2 is not an ansatz; its Osterwalder–Schrader provenance is machine-checked in the same tree, and we present it here because the reconstruction claim of this paper rests on it. All items below are green, in-tree prior work of the lane (modules `SpatialReflection`, `SpatialOS`, `SpatialReconstruction`); nothing in this subsection is new.

- *Reflection positivity over the measure.* The bond OS pairing of two half-chain observables is literally a sum over whole Gibbs paths — `osPairingBond_eq_gibbsSum`, through the splitting equivalence `bondEquiv` and the weight identity `gibbsWeight_joinBond` — and it is nonnegative on the diagonal for $\beta \geq 0$ (`osPairingBond_nonneg`; the restriction is sharp: `gibbsPathSum_neg_of_neg_beta` exhibits negativity at $\beta < 0$).
- *The collapse and its null space.* Half-chain observables collapse linearly onto boundary vectors (`collapseL`, surjective by `collapseL_surjective`), and the null space of the reflected site form is *exactly* the kernel of the collapse: $F \in \ker(\text{collapseL}) \iff \text{osPairingSite } F = 0$ (`mem_ker_collapseL_iff`). The quotient is by something identified, not by an abstract null space that happens to be there.
- *The physical quotient.* The quotient of half-chain observables by that kernel *is* the boundary space, as a bundled linear equivalence (`physicalEquiv`, via `quotKerEquivOfSurjective`).
- *The operator’s provenance.* The transfer operator’s matrix element between two collapsed half-chain observables is the reflected two-point sum of the Gibbs measure over whole paths, one slice further apart — `osPairing_transfer_gibbsSum`, stated for two observables, not one. Together with forcing (`transferOp_unique`), self-adjointness at every β (`transferOp_selfAdjoint`) and positivity for $\beta \geq 0$ (`transferOp_nonneg`), this is the OS construction of the operator whose spectrum Sections 4–7 control.

At finite extent the site form with $w > 0$ is definite, so on boundary vectors no further quotient intervenes; the quotient content lives one level up, on half-chain observables, where it is the `physicalEquiv` above. This is the honest division of labour between construction and spectrum in this model.

3 The dressing unitary

All statements of this and the following sections live in one module (`YangMills/OS/OSReconstructionUniform.lean` sha256 `d9919e256ae59ac9...27bac149`, 20,428 bytes, ten-pass ladder on the plane, final `ELAB_EXIT` 0). Write $Qu := \sqrt{w}u$ for the boundary dressing of a real Euclidean vector (`qEmbed`; it coincides with the measure side's `dress`, complexified — `qEmbed_eq_dress` is definitional); the site form weighs by $1/w$.

Theorem 3.1 (`A1, siteForm_qEmbed`). *For $w > 0$ and real u, v : $\text{siteForm } w (Qu) (Qv) = \sum_{\sigma} u_{\sigma} v_{\sigma}$. Per site the cancellation is elementary: $\sqrt{w}u \cdot \sqrt{w}v \cdot (1/w) = uv$.*

Theorem 3.2 (`A2, transferOp_qEmbed`). *$T(Qu) = Q(\text{act}(\text{symWeighted } w \beta) u)$ pointwise: the forced transfer operator on a dressed vector is the dressing of the symmetrised kernel's action.*

Theorem 3.3 (`A3, transferOp_iterate_qEmbed`). *$T^{[n]}(Qu) = Q((\text{act } S)^{[n]}u)$ — the shape the in-tree measure bridge `gibbsPathSum_eq_iterate` consumes.*

Together with `A1`, the dressing is a unitary between the Euclidean ℓ^2 pairing and the site form at every finite extent (the site form with $w > 0$ is definite, so no quotient intervenes — this is the elementary, finite-volume case of the OS construction, and we say so plainly).

4 The transported volume-uniform gap

Theorem 4.1 (`os_reconstruction_uniform_gap`). *Let $\beta, \gamma \in \mathbb{R}$ and $0 < \alpha < 1$ with $2 \tanh |\beta| + 2 \tanh |\gamma| \leq \alpha$. There is $m > 0$ such that for every L there are $\lambda > 0$ and $\Omega > 0$ with the fixed-point equation $\sum_{\tau} \text{tiltKernel}(\text{sliceW } \gamma L) \beta \lambda \sigma \tau \Omega_{\tau} = \Omega_{\sigma}$, the projected bound $\|\text{projectedTransfer}(\text{opOf}(\text{tiltKernel } w \beta \lambda))\| \leq e^{-m}$, and the `A2` identity at $w := \text{sliceW } \gamma L$.*

The proof is an obtain on `dobrushin_ising_uniform_gap` followed by `A2`: *the same witnesses*, no new analytic content — which is precisely the point. The uniformity is real because it is D-6's; the reconstruction contributes the identification, not a new estimate.

5 Exact identification against the Gibbs sums

The bridge from the raw measure to the operator is a chain of *identities* — no bound appears until the gap is invoked, and none of the statements below needs the window: they hold at every real β , every $\lambda \neq 0$, every positive weight.

Lemma 5.1 (`B0b, act_symWeighted_eq_smul_act_tilt`). *$\text{act } S u = \lambda \cdot \text{act}(\text{tilt}) u$ for $\lambda \neq 0$, and iterating (`act_symWeighted_iterate_eq_smul_tilt`, via the scalar transport lemmas `act_smul_fun`, `act_iterate_smul_fun`): $(\text{act } S)^{[N]}u = \lambda^N \cdot (\text{act tilt})^{[N]}u$.*

Lemma 5.2 (`B0a, act_iterate_eq_opOf_pow`). *$(\text{act } M)^{[n]}u = (\text{opOf } M)^n(\text{toLp } u)$ pointwise, for any kernel matrix M .*

Theorem 5.3 (`B1, gibbsPathSum_eq_inner_pow`). *For $w > 0$, $\beta \in \mathbb{R}$, $\lambda \neq 0$, $N \in \mathbb{N}$ and real observables A, B :*

$$\text{gibbsPathSum } w \beta N A B = \lambda^N \langle (\text{opOf}(\text{tiltKernel } w \beta \lambda))^N (QA), QB \rangle,$$

an exact identity: the unnormalised Gibbs two-point sum is λ^N times a matrix element of the reconstructed operator's powers, taken between dressed observables.

Theorem 5.4 (`B2`, `gibbsPartition_eq_inner_pow`). *Under the same hypotheses, the partition function is the same shape at the dressed constant: $\text{gibbsPartition } w \beta N = \lambda^N \langle \hat{T}^N(Q\mathbf{1}), Q\mathbf{1} \rangle$.*

Remark 5.5. *Dividing B1 by B2 cancels λ^N , so the normalised two-point function equals $\langle \hat{T}^N QA, QB \rangle / \langle \hat{T}^N Q\mathbf{1}, Q\mathbf{1} \rangle$ — plain arithmetic from the two identities, recorded here as prose. The quantitative six-term connected bound is a theorem of the same module, in division-free form: Section 8. We print the boundary between identity and estimate exactly where the kernel puts it.*

6 Mixed clustering on a generic Hilbert space

The clustering machinery is stated for an arbitrary real Hilbert space H (a `section` with `[NormedAddCommGroup H] [InnerProductSpace ℝ H] [CompleteSpace H]`), so the three lemmas below are reusable against any transfer operator packaged as a `VacuumTransfer` — in particular they are exactly the consumer shape of the `GappedTransfer` hypothesis of the abstract frameworks in Section 1.1.

Lemma 6.1 (`norm_sub_inner_smul_le`). *For $\|\Omega\| = 1$: $\|v - \langle \Omega, v \rangle \Omega\| \leq \|v\|$ (Pythagoras; removing the vacuum component does not grow the norm).*

Lemma 6.2 (`pow_apply_norm_le`). *$\|S^n v\| \leq r^n \|v\|$ whenever $\|S\| \leq r$, $r \geq 0$.*

Theorem 6.3 (`mixed_connCorr_bound`). *Let T be a `VacuumTransfer` with vacuum Ω , and suppose $\|\text{projectedTransfer } T \Omega\| \leq r$ with $r \geq 0$. Then for all $u, v \in H$ and all $n \in \mathbb{N}$ including $n = 0$:*

$$|\langle u, T^n v \rangle - \langle \Omega, u \rangle \langle \Omega, v \rangle| \leq \|u\| \|v\| r^n.$$

The $n = 0$ case is the Pythagoras lemma; $n \geq 1$ goes through the in-tree `projected_pow_succ`.

This strictly extends the lane’s earlier diagonal clustering (`os_reconstruction_uniform_clustering`, $u = v$, constant $\|v\|^2$), which remains in the module: two *different* observables, and the time-zero correlator is bounded too — the zero-time case is where a naive geometric-series argument is silent, and it is carried by the projection geometry instead.

7 The endpoint

Theorem 7.1 (`os_reconstruction_measure_uniform`). *Let $\beta, \gamma \in \mathbb{R}$ and $0 < \alpha < 1$ with $2 \tanh |\beta| + 2 \tanh |\gamma| \leq \alpha$. There is one $m > 0$ such that for every spatial extent L there are $\lambda > 0$ and $\Omega > 0$ with:*

- (i) *the fixed-point equation for `tiltKernel (sliceW γ L) β λ` at Ω ;*
- (ii) *exact identification: for every N and all real observables A, B , $\text{gibbsPathSum (sliceW γ L) β } N A B = \lambda^N \langle \hat{T}^N(QA), QB \rangle$ with $\hat{T} = \text{opOf (tiltKernel } \dots)$;*
- (iii) *volume-uniform mixed decay: for all u, v in the slice Hilbert space and all $n \geq 0$, $|\langle u, \hat{T}^n v \rangle - \langle \text{vacOf } \Omega, u \rangle \langle \text{vacOf } \Omega, v \rangle| \leq \|u\| \|v\| (e^{-m})^n$.*

The witnesses of (i)–(iii) are `dobrushin_ising_uniform_gap`’s, consumed verbatim; the packaging is `vacuumTransfer_opOf + tiltKernel_symm`; (ii) is B1 at $w = \text{sliceW } \gamma L$ (whose positivity is the in-tree `sliceW_pos`); (iii) is Theorem `mixed_connCorr_bound` at $r = e^{-m}$. One theorem, three clauses, no new constant: the mass that clause (iii) exhibits is the mass D-6 proved, now measured against the raw Gibbs sums through the exact identity of clause (ii).

8 The division-safe six-term connected bound

The one step the previous version named open is now a theorem, in two layers. The generic layer lives in the same arbitrary-Hilbert-space section as the mixed bound:

Theorem 8.1 (`six_term_connected_bound`). *Let T be a `VacuumTransfer` with vacuum Ω , let $0 \leq r \leq 1$ with $\|\text{projectedTransfer } T \Omega\| \leq r$, and write $F(u, v) := \langle u, T^N v \rangle$. Then for all q_A, q_B, q_1 and all N :*

$$|F(q_1, q_1) F(q_A, q_B) - F(q_A, q_1) F(q_1, q_B)| \leq 6 \|q_A\| \|q_B\| \|q_1\|^2 r^N.$$

The proof is a ring identity in E -terms ($E(u, v) := F(u, v) - \langle \Omega, u \rangle \langle \Omega, v \rangle$): the pure product terms cancel *exactly*, six terms survive, each carrying at least one E factor bounded by Theorem `mixed_connCorr_bound`, and the two $E \cdot E$ terms spend one factor through $r^N \leq 1$. The constant 6 is the honest count of surviving terms, not an optimised one.

The model layer consumes it at the dressed observables, *in raw Gibbs terms*:

Theorem 8.2 (`os_reconstruction_connected_uniform`). *Let $\beta, \gamma \in \mathbb{R}$ and $0 < \alpha < 1$ with $2 \tanh |\beta| + 2 \tanh |\gamma| \leq \alpha$. There is one $m > 0$ such that for every spatial extent L there are $\lambda > 0, \Omega > 0$ with the fixed-point equation, and for every time depth N and all real observables A, B :*

$$|Z_N S_N(A, B) - S_N(A, \mathbf{1}) S_N(\mathbf{1}, B)| \leq \lambda^{2N} (6 \|QA\| \|QB\| \|Q\mathbf{1}\|^2) (e^{-m})^N,$$

where $S_N = \text{gibbsPathSum}$, $Z_N = \text{gibbsPartition}$, and $Q \cdot = \text{toLp}(\text{dress} \cdot)$.

This is the normalised bound multiplied through by the positive scale of Z_N^2 . To divide it one needs a lower bound for $\langle \hat{T}^N Q\mathbf{1}, Q\mathbf{1} \rangle$ that does not degrade with N , and that is the content of the next section.

9 The positive cone: a denominator floor uniform in N

A spectral argument gives the floor only along the even powers: writing $Q\mathbf{1} = \langle \Omega, Q\mathbf{1} \rangle \Omega + q^\perp$ one has $\langle \hat{T}^N Q\mathbf{1}, Q\mathbf{1} \rangle = \langle \Omega, Q\mathbf{1} \rangle^2 + \langle S^N q^\perp, q^\perp \rangle$ with S the projected operator, and the correction is sign-controlled only for N even (S^{2k} is positive semidefinite; S^{2k+1} is not). The obstruction is an artefact of the method, not of the object: the kernel here is entrywise *positive*, so the order-theoretic route closes at every N .

Lemma 9.1 (`tiltKernel_pos`). *For $w > 0$ and $\lambda > 0$, $\text{tiltKernel } w \beta \lambda$ is entrywise positive — at every real β , since `spatialKernel` is a product of positive bond factors (`spatialKernel_pos`). The vacuum is positive too (`vacOf_pos`).*

Lemma 9.2 (`cone propagation, opOf_pow_ge_smul_fix`). *Let M be entrywise nonnegative with $\text{opOf } M \Omega = \Omega$. If $c \Omega \leq v$ entrywise then $c \Omega \leq T^n v$ entrywise, for every n . (Induction on n ; the step is $\sum_\tau M_{\sigma\tau}(c \Omega_\tau) \leq \sum_\tau M_{\sigma\tau}(T^n v)_\tau$ and the fixed-point equation. No parity, no spectral input.)*

Theorem 9.3 (`the floor, inner_pow_floor`). *Let M be entrywise nonnegative, Ω the normalised positive fixed vector, v entrywise positive. Then there is $f > 0$ with $f \leq \langle T^n v, v \rangle$ for every n . Explicitly $f = c^2 \|\Omega\|^2$ where $c = \min_\sigma (v_\sigma / \Omega_\sigma) > 0$ is the in-tree finite minimum `minWeight`.*

Theorem 9.4 (`os_reconstruction_normalised_clustering`). *In the window there is one $m > 0$ such that for every spatial extent L : the partition function is positive at every depth, and there is a constant $C(A, B)$ — depending on the observables and on L , but not on the time depth N — with*

$$\left| \frac{S_N(A, B)}{Z_N} - \frac{S_N(A, \mathbf{1})}{Z_N} \frac{S_N(\mathbf{1}, B)}{Z_N} \right| \leq C(A, B) (e^{-m})^N, \quad C(A, B) = \frac{6 \|QA\| \|QB\| \|Q\mathbf{1}\|^2}{f^2}.$$

The three quotients are the expectations of the normalised Gibbs measure: $S_N(A, B)/Z_N = \langle A(X_0)B(X_N) \rangle$, $S_N(A, \mathbf{1})/Z_N = \langle A(X_0) \rangle$ and $S_N(\mathbf{1}, B)/Z_N = \langle B(X_N) \rangle$. So the statement is that the connected two-point function *of the measure* decays at the volume-uniform rate, with a time-independent prefactor. This is the step the earlier version of this manuscript named as open; it is now a theorem, and the honest residue is stated in Section 11: C still depends on L .

10 Reproduction

All numbers below are read from committed run artefacts; predictions were committed *before* their measurements (the programme’s standing discipline — the ledger records nine consecutive exact count predictions for this lane before this paper’s two).

Stage OS-R-1 (measured, ledger Addenda 616–619)

Quantity	Predicted	Measured
core jobs, unwired	8480 (ledger baseline)	8480
core jobs, wired	$8480 + 1$	8481
oracle reports, wired	$3010 + 7$	3017
<code>sorryAx</code> occurrences	0	0
endpoint axiom sets	standard triple	exactly [<code>propext</code> , <code>Classical.choice</code> , <code>Quot.sound</code>]

Fresh-clone witness at anchor `8127e5b52d75116b19c172ea022e12c1cab0f8c2`: judges 110/110 in normal and -0; core 8481; oracle 3017; module sha256 bit-identical (Addendum 619; run cell `scripts/osr1_pass7_freshclone.py`).

Stage OS-R-2 (predictions Addendum 621, committed before the run; measurements Addendum 622)

Quantity	Predicted (committed first)	Measured
core jobs	8481 exact (file grew, no module added)	8481
oracle reports	3017 + 9 = 3026 exact	3026
sorryAx occurrences	0	0
new endpoint axiom sets	standard triple	exactly [propext , Classical.choice ,
<i>six-term stage (predictions Addendum 623, committed before the run; measured Addendum 624):</i>		
core jobs	8481 exact (file grew again)	8481
oracle reports	3026 + 2 = 3028 exact	3028
sorryAx occurrences	0	0
new endpoint axiom sets	standard triple	exactly the triple $\times 2$
fresh-clone reproduction	bit-identical module, same counts	judges 110/110 both modes; 8481; 3028; 0
<i>positive-cone stage (predictions Addendum 626, committed before the run; measured Addendum 627):</i>		
core jobs	8481 exact (file grew a third time)	8481
oracle reports	3028 + 7 = 3035 exact	3035
sorryAx occurrences	0	0
new endpoint axiom sets	standard triple	exactly the triple $\times 7$
fresh-clone reproduction	bit-identical module, same counts	judges 110/110 both modes; 8481; 3035; 0

The core count is the tenth consecutive exact count prediction of this lane; the six-term stage made it eleven and the cone stage twelve: the module file grew from 12,347 to 20,428 to 32,533 to 43,213 bytes across the three stages and the job total never moved. Module: `YangMills/OS/OSReconstructionUnit` final sha256 `2f2f4cf95eb9a5ccb51f4b93ee7cf87b8cb49b34a3fe51eb746c34e4996e9e4e` (43,213 bytes), banked at commit `6ef9aee03` (six-term stage: `3b3598eb...`, 32,533 B at `b47dda50e`; measure stage: `d9919e25...`, 20,428 B at `3b0dfa9e8`) after the plane printed `ELAB_EXIT 0` on byte-identical source (the cell prints size and sha256 before elaborating; both matched the local file before banking). Pass history: passes 1–5 (OS-R-1, Addenda 614–616), passes 8–10 (OS-R-2, Addendum 621: the inner-product notation trap at this Mathlib pin, then a single `nlinarith` orientation failure in the Pythagoras lemma, then green). The wiring unit with the registered predictions is committed as `scripts/osr2_pass11_wiring.py` *before* its run, at commit `520e177a7`.

11 What is not claimed

No Wightman axioms, no $OS \rightarrow$ Wightman analytic continuation, no continuum limit, no $SU(N)$, no statement about the Clay problem. The window is the Dobrushin window; outside it the programme’s own kill-test (repository, Addendum 610) shows the standard-cone Birkhoff route is projectively blind to the site weight, which is why no claim beyond the window appears here. The model is the $\mathbb{Z}_2$ Ising slice; “reconstruction” means the elementary finite-volume site-form construction (the quotient content lives on half-chain observables, Section 2.1), and the theorems are about its *spectrum and its exact relation to the Gibbs sums*, not about the construction being difficult. The connected bound is printed in both forms, division-free (Section 8) and normalised (Section 9); the residue there is precise and we state it rather than blur it: the mass m is uniform in the spatial extent L , but the *constant* $C(A, B)$ of the normalised bound is not — it carries $1/f^2$ with $f = c^2 \|\Omega\|^2$ and c the minimum ratio $Q1/\Omega$, which may degrade as $L \rightarrow \infty$. So this paper proves a volume-uniform *rate* with per-extent constants, not a volume-uniform bound; making C uniform in L as well is the next step, and it is a genuinely new estimate rather than a further identification.

References

- [1] M. R. Douglas, S. Hoback, C. Mei, Y. Nissim, *Formalization of QFT*, arXiv:2603.15770 (2026). Repositories: `mrddouglasny/OSforGFF`, `mrddouglasny/reflection-positivity`, `mrddouglasny/lgt`, `mrddouglasny/pphi2`; `xiyin137/OSreconstruction`.
- [2] K. Osterwalder, R. Schrader, *Axioms for Euclidean Green's functions* I, II, *Comm. Math. Phys.* **31** (1973), **42** (1975).
- [3] J. Glimm, A. Jaffe, *Quantum Physics: A Functional Integral Point of View*, 2nd ed., Springer (1987).
- [4] R. L. Dobrushin, *The description of a random field by means of conditional probabilities and conditions of its regularity*, *Theory Probab. Appl.* **13** (1968).
- [5] L. Eriksson, *From the Gibbs Weight to the Spectral Gap: A Complete Machine-Checked Osterwalder–Seiler Chain for the $\mathbb{Z}_2$ Lattice Gauge Chain*, ai.viXra:2607.0078 (2026) — the zero-extension, explicitly non-volume-uniform predecessor of this paper.
- [6] L. Eriksson, *the dobrushin-matrix paper of the same repository* (`papers/dobrushin-matrix/`), source of `dobrushin_ising_uniform_gap` (D-6, ledger Addendum 603).