

The Rate Without the Extent: a Machine-Checked Extent-Independent Spectral Rate for the Decoupled $\mathbb{Z}_2$ Slice and a pre-registered gate we failed to respect

Lluis Eriksson

Abstract

Every rate in this lane so far has been a fixed-extent rate. The gap paper proved strict spectral separation at each extent and said plainly that it was not uniform [2]; the modulus paper gave that separation a number, $\text{specRatio}(L)$, and reported measurements saying the number tends to 1 outside the disordered region [3]. A geometric bound whose rate tends to 1 is empty in the volume limit, so nothing in the lane survived $L \rightarrow \infty$, and the word *clustering* was never used.

Proved, for $\beta \geq 0$ throughout. For the *decoupled* kernel — the transfer kernel at constant source weight — the modulus is $\text{specRatio} = \tanh \beta$ at every *positive* extent, with L nowhere in it. (At $L = 0$ there is a single configuration, the only eigenvalue is the Perron one and $\text{specRatio} = 0$. The exact identity is stated for $L \geq 1$; the operator and Gibbs inequalities below also hold at $L = 0$, where they are trivial.) Both directions: an operator bound by induction on the extent, and attainment by the single-site observable the extent paper already built [1]. The hypothesis $\beta \geq 0$ is load-bearing, not a convention: at negative coupling the odd bond eigenvalue is negative and $\tanh \beta \cdot Z^L$ is not specGap .

And a bound that is uniform, not only a rate that is. An extent-independent rate settles nothing on its own: $C(L) \rho^N$ with $C(L) \rightarrow \infty$ is empty as the system grows however small ρ is. At constant source weight everything is explicit, so the uniform bound is computed rather than composed: for *every* N , with no threshold,

$$|\mathbb{E}[A(X_0)A(X_N)]| \leq \langle A^2 \rangle (\tanh \beta)^N,$$

with $\langle A^2 \rangle$ the mean square in the uniform measure — at most 1 for $|A| \leq 1$, at every extent. We do *not* call this clustering: no infinite-volume state is constructed here, and none is claimed.

Why the proof is not the spectral decomposition. The decoupled kernel is a product over sites, so its spectrum is a product; that route needs the spectrum of a Kronecker power, which the library does not carry. It is not needed. The modulus paper proved that specGap is the *greatest* norm ratio on the fluctuation sector, so bounding it above is an operator inequality and nothing else, and the operator inequality falls to induction: the even part of an observable keeps its mean zero and inherits the rate, the odd part keeps nothing and gets only Schur’s test, and the two recombine *exactly*, because $\tanh \beta \cdot Z = D$ with Z the row sum and D the odd eigenvalue of a single bond.

Not proved, and a gate we failed to respect. The *coupled* kernel is untouched. Before any of this was written we pre-registered two falsifiable predictions. The first — the decoupled rate is exactly $\tanh \beta$ at every extent — passed to 10^{-16} . The second — the coupled uniformity boundary is the Onsager curve — failed on one of eight cells. **The script we committed had a single joint gate, and with the second prediction failing that**

gate read FAIL -- do not fabricate. We fabricated anyway. The theorems below do not depend on either prediction and stand on their proofs; the *process* claim does not, and §6 records the violation rather than reinterpreting the gate after the fact. The Onsager boundary claim is reported as **not established**. At constant source weight the Gibbs measure factorises across spatial *sites* into L independent one-dimensional temporal chains — the slices $X_0, \dots, X_N$ themselves stay coupled in time, which is what the rate $\tanh \beta$ measures — so what is proved here is a statement about a product measure over sites; that is exactly why it is reachable, and it is said here rather than left to be noticed. Reflection positivity is untouched, and nothing in this paper is a claim about $SU(N)$, the continuum limit, or the Yang–Mills mass gap [10].

1 Ten papers of fixed-extent rates

The lane’s rates have all been local in the extent. That is not a stylistic complaint: a rate $r(L)$ with $r(L) \rightarrow 1$ gives, at fixed N , a bound that becomes vacuous as the system grows, and the infinite-volume statement one actually wants — exponential decay of correlations with a fixed rate — is not implied by any amount of fixed- L information.

So the honest question is whether *some* kernel in this family has a rate that does not move with L . This paper answers it for the decoupled kernel, and the answer is exact.

2 The construction

Standing hypothesis: $\beta \geq 0$. Every statement below carries it, and it is load-bearing. At $\beta < 0$ the number D is negative while `specGap` is by definition non-negative, so Theorem 4.1 would assert that a non-negative quantity is negative. (The all- β form would read $|\tanh \beta|$; we simply assume $\beta \geq 0$, as the formalisation does.)

Let $Z = e^\beta + e^{-\beta}$ and $D = e^\beta - e^{-\beta}$, and let K be the decoupled kernel on $\{0, 1\}^L$: a product over sites of the single-bond matrix. Its row sums are all Z^L , so the constant observable is its Perron vector with eigenvalue $\lambda = Z^L$.

Lemma 2.1 (Schur’s test). *A symmetric nonnegative kernel all of whose row sums equal R satisfies $\|Ku\| \leq R \|u\|$.*

`norm_act_le_of_rowSum`. One Cauchy–Schwarz per row, then summing. This is the only bound the odd part of the induction is given, and it is all it needs.

Theorem 2.2 (the uniform operator bound). *Let $\beta \geq 0$. For every extent L and every observable u with $\sum_\sigma u(\sigma) = 0$,*

$$\|Ku\| \leq \tanh \beta \cdot Z^L \cdot \|u\|.$$

`spatialKernel_fluct_bound`, from the squared form `spatialKernel_fluct_sq`, which is what the induction runs on — a square root in the inductive hypothesis would be a square root at every step.

3 Why the constant does not degrade

Split an observable at the first site, $u = (u_0, u_1)$, and set $p = u_0 + u_1$, $m = u_0 - u_1$. Then p still has mean zero, so the inductive hypothesis applies to it with rate $\tanh \beta$; m has no orthogonality

left and gets only Lemma 2.1, with rate 1. Acting with the kernel and collecting,

$$\|Ku\|^2 = \frac{1}{2}(Z^2\|Kp\|^2 + D^2\|Km\|^2), \quad \|u\|^2 = \frac{1}{2}(\|p\|^2 + \|m\|^2).$$

Feeding in the two bounds gives $\frac{1}{2}Z^2(\tanh \beta Z^L)^2\|p\|^2 + \frac{1}{2}D^2(Z^L)^2\|m\|^2$, and the two coefficients are *equal*, because

$$\tanh \beta \cdot Z = D.$$

tanh_mul_z2Norm. That single identity is the whole reason the constant survives: the rate given up on the odd part is exactly the odd eigenvalue of one bond, at every step. Lose it and the constant would pick up a factor per site.

4 Sharpness

Theorem 2.2 is an inequality. It is attained.

Theorem 4.1 (the rate is exactly $\tanh \beta$). *Let $\beta \geq 0$. At every extent with at least one site, $\text{specGap} = \tanh \beta \cdot Z^L$, hence $\text{specRatio} = \tanh \beta$.*

spatialKernel_specGap_eq. The upper bound is Theorem 2.2 together with the fact that specGap is the *greatest* norm ratio on the fluctuation sector [3] — which is the only *spectral* use of that paper here, and it is used as a bound, never as a decomposition. (§5 uses that paper again, for its Gibbs endpoint.) The lower bound is the single-site sign observable of the extent paper [1]: it has mean zero (**sum_siteSign**), it is nonzero, and its eigenvalue $D Z^{L-1}$ realises the bound exactly.

5 The payoff

Theorem 5.1 (an extent-independent rate). *Let $\beta \geq 0$. At constant source weight there is an N_0 such that for every observable A whose dressing has mean zero there is a $C_A > 0$ with $|\mathbb{E}[A(X_0)A(X_N)]| \leq C_A (\tanh \beta)^N$ for $N \geq N_0$, and $\tanh \beta$ does not depend on the extent.*

gibbs_decay_extent_free_rate, composing Theorem 4.1 with the modulus paper's endpoint [3].

This is not a volume limit, and the previous version said it was. In Theorem 5.1 the extent is fixed before N_0 and C_A are chosen, so both may in principle diverge with L , and a bound $C(L)\rho^N$ with $C(L) \rightarrow \infty$ says nothing as the system grows however small ρ is. An extent-independent rate is a necessary ingredient of an infinite-volume statement, not a sufficient one.

Theorem 5.2 (the bound, and not merely the rate, is uniform). *Let $\beta \geq 0$. At constant source weight, for every observable A with $\sum_\sigma A(\sigma) = 0$ and every N — no threshold —*

$$|\mathbb{E}[A(X_0)A(X_N)]| \leq \langle A^2 \rangle (\tanh \beta)^N, \quad \langle A^2 \rangle = 2^{-L} \sum_\sigma A(\sigma)^2.$$

gibbsCorr_one_uniform_bound. Here nothing is left to depend on the extent except through the observable itself: for $|A| \leq 1$ the constant is at most 1 at every L . It is computed rather than composed — at constant weight the dressed constant observable *is* the Perron vector and

the partition function is exactly $2^L Z^{LN}$, so the endpoint’s threshold machinery is not needed at all.

Two immediate consequences are worth stating as theorems rather than left to the reader, since each removes a hypothesis or a symbol.

Corollary 5.3 (arbitrary observable, in covariance form). *Let $\beta \geq 0$. For every observable A , centring in the uniform measure gives $|\mathbb{E}[\bar{A}(X_0)\bar{A}(X_N)]| \leq \text{Var}_{\text{unif}}(A) (\tanh \beta)^N$, where $\bar{A} = A - \langle A \rangle_{\text{unif}}$.*

[gibbsCov_one_uniform_bound](#).

Corollary 5.4 (no L on the right at all). *Let $\beta \geq 0$. If A has mean zero and $|A| \leq 1$ pointwise, then for every N ,*

$$|\mathbb{E}[A(X_0)A(X_N)]| \leq (\tanh \beta)^N.$$

[gibbsCorr_one_le_of_bounded](#). Every symbol on the right is free of the extent, and there is no threshold and no constant. This is the sharpest form the product case admits.

We still do not call this clustering. No infinite-volume state is constructed in this paper, and a uniform finite-volume bound for a fixed sequence of observables is not by itself such a construction.

6 Scope, and a judge that failed

This is a product measure. At constant source weight the spatial slices decouple across *sites*: the Gibbs measure factorises into L independent one-dimensional temporal chains, and the statement above is, in substance, one such chain repeated. The slices $X_0, \dots, X_N$ are of course not independent of each other — their dependence is precisely what $\tanh \beta$ bounds. That is why it is reachable. What it establishes is narrower than it may sound: that *the extent alone does not degrade the rate* in the product kernel, so the unresolved dependence enters only once the spatial source weight is switched on. It does not show that every nonconstant weight destroys uniformity, nor that the weight is the sole cause of the coupled model’s degradation.

The coupled kernel is exactly a 2D Ising transfer matrix. With the ring source weight, $K = D_w^{1/2} T D_w^{1/2}$ where T is the inter-slice product and D_w the intra-slice ring weight — which is the symmetrised transfer matrix of the two-dimensional Ising model, vertical coupling β and horizontal coupling γ . This is an exact algebraic identification, not a numerical claim.

The pre-registered gate, and our failure to respect it. Before any of the above was written we committed two falsifiable predictions with their thresholds fixed in advance [scripts/judge_spatial_uniform.py](#). The first — the decoupled rate is exactly $\tanh \beta$ at every extent — passed to 10^{-16} . The second — that the *coupled* uniformity boundary is the Onsager curve $\sinh 2\beta \sinh 2\gamma = 1$ — predicted saturation below 1 on the four disordered cells and a rate tending to 1 on the four ordered ones. Seven of eight cells agreed.

But the script we committed ends in a single joint gate, printing `PASS -- fabrication authorised` only when *both* predictions hold. With the second failing, the artefact printed `FAIL -- do not fabricate`, and we proceeded to prove the decoupled theorem anyway. Deciding afterwards that the two predictions had separate gates is exactly the retroactive reinterpretation pre-registration exists to prevent, so we do not claim it: the gate was joint,

it read FAIL, and we did not respect it. The theorems of §§2–5 rest on their proofs and are unaffected; what is forfeited is the methodological claim, and the correct remedy — writing the two gates separately *before* looking — is available only to the next campaign, not to this one. The cell $(\beta, \gamma) = (0.1, 1.0)$, disordered, gave $\text{specRatio}(12) = 0.6266$ — nowhere near 1 — but with drift from $L = 10$ of 3.5×10^{-2} , above the pre-registered 10^{-2} , so the cell is *undecided at* $L \leq 12$ and the judge does not pass.

We did not raise the extent until the cell passed. The autopsy [scripts/judge_spatial_uniform_autopsy.py](#) prints the whole sequence instead: its last increment ratio is 0.835, and a geometric extrapolation from it gives about 0.707. That single local ratio is *consistent with* saturation below one; it does not exclude a later change of regime, and $L \leq 12$ does not establish a limiting value. **The Onsager identification of the boundary is reported as not established**, and no statement in this paper rests on it.

No claim about the physical target. Reflection positivity is untouched, and nothing here concerns $\text{SU}(N)$, the continuum limit, or the Clay problem [10]. The distance to that problem is unchanged.

7 What is verified and not proved

The measured $\text{specRatio}(L)$ for the coupled kernel, extended past the five points the modulus paper printed:

(β, γ)	$L = 1$	$L = 12$	increments
(0.2, 0.2)	0.1974	0.294449	geometric, ratio 0.278
(0.3, 0.4)	0.2913	0.646895	geometric, ratio 0.648
(0.1, 1.0)	0.0997	0.626639	geometric, ratio 0.835
(0.8, 1.2)	0.6640	1.000000	to 1

These are **verified** and **not proved**; no theorem depends on them. They are printed because they are the reason the coupled case is believed hard rather than false, and because the last row is the one the whole lane has been reporting as its limitation.

8 Reproducibility

All artifacts are at source checkpoint [2427eb49a36e755c480e6161b2f7beeeef2ea52](#) on branch `main`. Every hyperlink was resolved against that commit, and checked to point at an actual declaration line, before compilation.

Quantity	Value
Full core build	8431 jobs, success
Oracle commands, and answers	2697
distinct declarations	2694
commands listed twice	3
distinct, with axiom dependencies	2671
distinct, axiom-free	23
Declarations in this module	25
Distinct axioms across all outputs	<code>propext</code> , <code>Quot.sound</code> , <code>Classical.choice</code>
Occurrences of <code>sorryAx</code>	0
Project axioms	0
Errors	0

The module is [YangMills/OS/SpatialUniform.lean](#); the pre-registered judge and its autopsy are [scripts/judge_spatial_uniform.py](#) and [scripts/judge_spatial_uniform_autopsy.py](#); verdicts are in [docs/VERIFICATION-LEDGER.md](#).

References

- [1] L. Eriksson, *Where the Elementary Reconstruction Stops: a Machine-Checked Obstruction at the Coupled $\mathbb{Z}_2$ Slice*, viXra:2607.0075, 2026.
- [2] L. Eriksson, *Strict but Not Uniform: a Machine-Checked Spectral Gap at Every Finite Extent*, viXra, 2026.
- [3] L. Eriksson, *The Modulus: a Machine-Checked Operator Bound on the Fluctuation Sector of the Coupled $\mathbb{Z}_2$ Slice*, viXra, 2026.
- [4] L. Onsager, *Crystal statistics I: a two-dimensional model with an order-disorder transition*, Phys. Rev. **65** (1944), 117–149.
- [5] M. Kac and J. C. Ward, *A combinatorial solution of the two-dimensional Ising model*, Phys. Rev. **88** (1952), 1332–1337.
- [6] T. D. Schultz, D. C. Mattis and E. H. Lieb, *Two-dimensional Ising model as a soluble problem of many fermions*, Rev. Mod. Phys. **36** (1964), 856–871.
- [7] D. A. Levin and Y. Peres, *Markov Chains and Mixing Times*, 2nd ed., AMS, 2017. (Tensorisation of the spectral gap for product chains; Chapter 12.)
- [8] P. Diaconis and L. Saloff-Coste, *Comparison theorems for reversible Markov chains*, Ann. Appl. Probab. **3** (1993), 696–730.
- [9] E. Seneta, *Non-negative Matrices and Markov Chains*, 2nd ed., Springer, 1981.
- [10] A. Jaffe and E. Witten, *Quantum Yang–Mills theory*, Clay Mathematics Institute Millennium Prize problem description, 2000.